

Submitted by Navneet Das 3433 Comp A

```
***  
Mini Project 1 Deep Learning  
Face Recognition  
***
```

```
In [ ]: import numpy as np # Linear algebra  
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)  
  
import os  
for dirname, _, filenames in os.walk('/kaggle/input'):  
    for filename in filenames:  
        print(os.path.join(dirname, filename))  
  
/kaggle/input/opencv-facial-recognition-lbph/yalefaces/test/subject03.glasses.gif  
/kaggle/input/opencv-facial-recognition-lbph/yalefaces/test/subject12.normal.gif  
/kaggle/input/opencv-facial-recognition-lbph/yalefaces/test/subject02.leftlight.gif  
/kaggle/input/opencv-facial-recognition-lbph/yalefaces/test/subject13.sad.gif  
/kaggle/input/opencv-facial-recognition-lbph/yalefaces/test/subject06.leftlight.gif  
/kaggle/input/opencv-facial-recognition-lbph/yalefaces/test/subject11.glasses.gif  
/kaggle/input/opencv-facial-recognition-lbph/yalefaces/test/subject02.centerlight.gif  
/kaggle/input/opencv-facial-recognition-lbph/yalefaces/test/subject14.sad.gif  
/kaggle/input/opencv-facial-recognition-lbph/yalefaces/test/subject14.normal.gif  
/kaggle/input/opencv-facial-recognition-lbph/yalefaces/test/subject04.surprised.gif  
/kaggle/input/opencv-facial-recognition-lbph/yalefaces/test/subject01.happy.gif  
/kaggle/input/opencv-facial-recognition-lbph/yalefaces/test/subject09.rightlight.gif  
/kaggle/input/opencv-facial-recognition-lbph/yalefaces/test/subject15.rightlight.gif  
/kaggle/input/opencv-facial-recognition-lbph/yalefaces/test/subject09.sad.gif  
/kaggle/input/opencv-facial-recognition-lbph/yalefaces/test/subject15.sad.gif  
/kaggle/input/opencv-facial-recognition-lbph/yalefaces/test/subject07.happy.gif  
/kaggle/input/opencv-facial-recognition-lbph/yalefaces/test/subject12.rightlight.gif  
/kaggle/input/opencv-facial-recognition-lbph/yalefaces/test/subject05.surprised.gif  
/kaggle/input/opencv-facial-recognition-lbph/yalefaces/test/subject04.leftlight.gif
```

```
In [ ]: pip install mtcnn  
  
Collecting mtcnn  
  Downloading mtcnn-0.1.1-py3-none-any.whl (2.3 MB)  
    |████████████████████████████████████████| 2.3 MB 567 kB/s  
Requirement already satisfied: keras>=2.0.0 in /opt/conda/lib/python3.7/site-packages (from mtcnn) (2.6.0)  
Requirement already satisfied: opencv-python>=4.1.0 in /opt/conda/lib/python3.7/site-packages (from mtcnn) (4.5.4.60)  
Requirement already satisfied: numpy>=1.14.5 in /opt/conda/lib/python3.7/site-packages (from opencv-python>=4.1.0->mtcnn) (1.19.5)  
Installing collected packages: mtcnn  
Successfully installed mtcnn-0.1.1  
WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. It is recommended to use a virtual environment instead: https://pip.pypa.io/warnings/venv (https://pip.pypa.io/warnings/venv)  
Note: you may need to restart the kernel to use updated packages.
```

```
In [ ]: from matplotlib import pyplot as plt  
import cv2  
from PIL import Image  
  
import mtcnn  
from mtcnn.mtcnn import MTCNN  
from matplotlib.patches import Rectangle  
  
from os import listdir  
from tqdm import tqdm  
import pandas as pd  
from sklearn.model_selection import train_test_split  
import seaborn as sns
```

```
In [ ]: DIRECTORY = "../input/opencv-facial-recognition-lbph/yalefaces/"  
DIRECTORY_train='../input/opencv-facial-recognition-lbph/yalefaces/train/'  
DIRECTORY_test='../input/opencv-facial-recognition-lbph/yalefaces/test/'
```

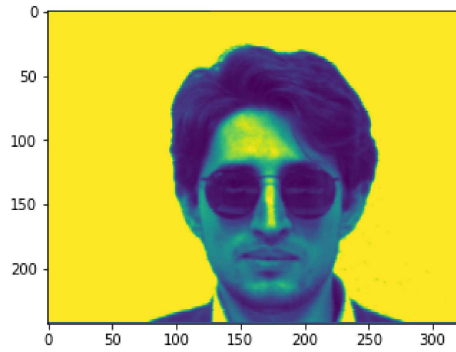
1.Data Analysis

1) Data Sample

```
In [ ]: filename = "../input/opencv-facial-recognition-lbph/yalefaces/test/subject03.glasses.gif"
pixels = plt.imread(filename)

rgb_pixels = np.stack((pixels, pixels, pixels), axis=2)
print(rgb_pixels.shape)
plt.imshow(pixels)
plt.show()
```

(243, 320, 3)



2) MTCNN

```
In [ ]: detector = MTCNN()  
        results = detector.detect_faces(rgb_pixels)  
        results
```

User settings:

```
KMP_AFFINITY=granularity=fine,verbose,compact,1,0
KMP_BLOCKTIME=0
KMP_DUPLICATE_LIB_OK=True
KMP_INIT_AT_FORK=FALSE
KMP_SETTINGS=1
KMP_WARNINGS=0
```

Effective settings:

```
KMP_ABORT_DELAY=0
KMP_ADAPTIVE_LOCK_PROPS='1,1024'
KMP_ALIGN_ALLOC=64
KMP_ALL_THREADPRIVATE=128
KMP_ATOMIC_MODE=2
KMP_BLOCKTIME=0
KMP_CPUINFO_FILE: value is not defined
KMP_DETERMINISTIC_REDUCTION=false
KMP_DEVICE_THREAD_LIMIT=2147483647
KMP_DISP_NUM_BUFFERS=7
KMP_DUPLICATE_LIB_OK=true
KMP_ENABLE_TASK_THROTTLING=true
KMP_FORCE_REDUCTION: value is not defined
KMP_FOREIGN_THREADS_THREADPRIVATE=true
KMP_FORKJOIN_BARRIER='2,2'
KMP_FORKJOIN_BARRIER_PATTERN='hyper,hyper'
KMP_GTID_MODE=3
KMP_HANDLE_SIGNALS=false
KMP_HOT_TEAMS_MAX_LEVEL=1
KMP_HOT_TEAMS_MODE=0
KMP_INIT_AT_FORK=true
KMP_LIBRARY=throughput
KMP_LOCK_KIND=queuing
KMP_MALLOC_POOL_INCR=1M
KMP_NUM_LOCKS_IN_BLOCK=1
KMP_PLAIN_BARRIER='2,2'
KMP_PLAIN_BARRIER_PATTERN='hyper,hyper'
KMP_REDUCTION_BARRIER='1,1'
KMP_REDUCTION_BARRIER_PATTERN='hyper,hyper'
KMP_SCHEDULE='static,balanced;guided,iterative'
KMP_SETTINGS=true
KMP_SPIN_BACKOFF_PARAMS='4096,100'
KMP_STACKOFFSET=64
KMP_STACKPAD=0
KMP_STACKSIZE=8M
KMP_STORAGE_MAP=false
KMP_TASKING=2
KMP_TASKLOOP_MIN_TASKS=0
KMP_TASK_STEALING_CONSTRAINT=1
KMP_TEAMS_THREAD_LIMIT=4
KMP_TOPOLOGY_METHOD=all
KMP_USE_YIELD=1
KMP_VERSION=false
KMP_WARNINGS=false
OMP_AFFINITY_FORMAT='OMP: pid %P tid %i thread %n bound to OS proc set {%A}'
OMP_ALLOCATOR=omp_default_mem_alloc
OMP_CANCELLATION=false
OMP_DEFAULT_DEVICE=0
OMP_DISPLAY_AFFINITY=false
OMP_DISPLAY_ENV=false
OMP_DYNAMIC=false
OMP_MAX_ACTIVE_LEVELS=1
OMP_MAX_TASK_PRIORITY=0
OMP_NESTED: deprecated; max-active-levels-var=1
OMP_NUM_THREADS: value is not defined
OMP_PLACES: value is not defined
OMP_PROC_BIND='intel'
OMP_SCHEDULE='static'
OMP_STACKSIZE=8M
OMP_TARGET_OFFLOAD=DEFAULT
OMP_THREAD_LIMIT=2147483647
OMP_WAIT_POLICY=PASSIVE
KMP_AFFINITY='verbose,warnings,respect,granularity=fine,compact,1,0'
```

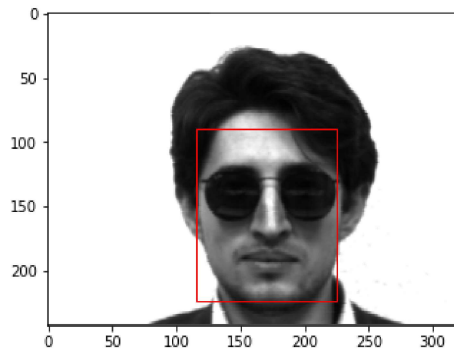
```
2021-12-06 14:52:23.044469: I tensorflow/core/common_runtime/process_util.cc:146] Creating new thread pool with default inter op setting: 2. Tune using inter_op_parallelism_threads for best performance.
2021-12-06 14:52:23.417070: I tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:185] None of the MLIR Optimizat
```

ion Passes are enabled (registered 2)

```
Out[6]: [{'box': [115, 90, 110, 134],
'confidence': 0.9997938275337219,
'keypoints': {'left_eye': (147, 140),
'right_eye': (198, 138),
'nose': (174, 169),
'mouth_left': (151, 192),
'mouth_right': (195, 193)}]}
```

```
In [ ]: def draw_image_with_boxes(data, result_list):
plt.imshow(data)
ax = plt.gca()
for result in result_list:
    x, y, width, height = result['box']
    rect = Rectangle((x, y), width, height, fill=False, color='red')
    ax.add_patch(rect)
plt.show()

draw_image_with_boxes(rgb_pixels, results)
```



3) Extract and normalise the face pixels

```
In [ ]: def extract_face_from_file(filename, required_size=(160, 160)):
        image = Image.open(filename)

        return extract_face(image, required_size)

def extract_face(image, required_size=(160, 160)):
    image = image.convert('RGB')
    pixels = np.asarray(image)
    results = detector.detect_faces(pixels)

    x1, y1, width, height = results[0]['box']

    x1, y1 = abs(x1), abs(y1)
    x2, y2 = x1 + width, y1 + height

    face = pixels[y1:y2, x1:x2]

    image = Image.fromarray(face)
    image = image.resize(required_size)
    face_array = np.asarray(image)
    gray_face = cv2.cvtColor(face_array, cv2.COLOR_BGR2GRAY)

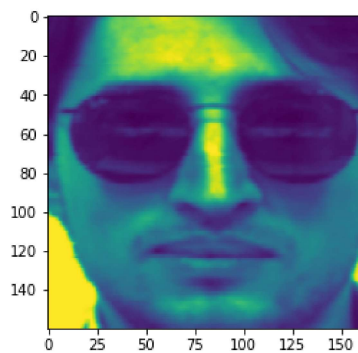
    return gray_face

detector = MTCNN()

face_pixels = extract_face_from_file("../input/opencv-facial-recognition-lbph/yalefaces/test/subject03.glasses.gif")

plt.imshow(face_pixels)
```

Out[8]: <matplotlib.image.AxesImage at 0x7f4e517775d0>




2.Data set ~ 135 train datas and 30 test datas

```
In [ ]: def list_files(directory, contains):
        return list(f for f in listdir(directory) if contains in f)
```

```
In [ ]: i = 1
faces = list()
for filename in tqdm(list_files(DIRECTORY_train, "subject")[0:16]):
    # path
    path = DIRECTORY_train + filename
    # get face
    face = extract_face_from_file(path)
    # plot
    plt.subplot(4, 4, i)
    plt.axis('off')
    plt.imshow(face)
    faces.append(face)
    i += 1
plt.show()
```

100%|██████████| 16/16 [00:09<00:00, 1.70it/s]



```
In [ ]: filenames = pd.DataFrame(list_files(DIRECTORY_train, "subject"))

df_train = filenames[0].str.split(".", expand=True)
df_train["filename"] = filenames

df_train = df_train.rename(columns = {0:"subject", 1:"category"})
df_train['subject'] = df_train.subject.str.replace('subject', '')
df_train.apply(pd.to_numeric, errors='coerce').dropna()
df_train['subject'] = pd.to_numeric(df_train["subject"])
df_train
```

```
Out[11]:
```

	subject	category	2	filename
0	15	happy	gif	subject15.happy.gif
1	13	noglasses	gif	subject13.noglasses.gif
2	1	normal	gif	subject01.normal.gif
3	15	surprised	gif	subject15.surprised.gif
4	14	wink	gif	subject14.wink.gif
...
130	5	sad	gif	subject05.sad.gif
131	3	surprised	gif	subject03.surprised.gif
132	4	sad	gif	subject04.sad.gif
133	7	normal	gif	subject07.normal.gif
134	2	surprised	gif	subject02.surprised.gif

135 rows × 4 columns

```

In [ ]: filenames2 = pd.DataFrame(list_files(DIRECTORY_test, "subject"))

df_test = filenames2[0].str.split(".", expand=True)
df_test["filename"] = filenames2

df_test = df_test.rename(columns = {0:"subject", 1:"category"})
df_test['subject'] = df_test.subject.str.replace('subject', '')
df_test.apply(pd.to_numeric, errors='coerce').dropna()
df_test['subject'] = pd.to_numeric(df_test["subject"])
df_test

```

```

Out[12]:

```

	subject	category	2	filename
0	3	glasses	gif	subject03.glasses.gif
1	12	normal	gif	subject12.normal.gif
2	2	leftlight	gif	subject02.leftlight.gif
3	13	sad	gif	subject13.sad.gif
4	6	leftlight	gif	subject06.leftlight.gif
5	11	glasses	gif	subject11.glasses.gif
6	2	centerlight	gif	subject02.centerlight.gif
7	14	sad	gif	subject14.sad.gif
8	14	normal	gif	subject14.normal.gif
9	4	surprised	gif	subject04.surprised.gif
10	1	happy	gif	subject01.happy.gif
11	9	rightlight	gif	subject09.rightlight.gif
12	15	rightlight	gif	subject15.rightlight.gif
13	9	sad	gif	subject09.sad.gif
14	15	sad	gif	subject15.sad.gif
15	7	happy	gif	subject07.happy.gif
16	12	rightlight	gif	subject12.rightlight.gif
17	5	surprised	gif	subject05.surprised.gif
18	4	leftlight	gif	subject04.leftlight.gif
19	10	sad	gif	subject10.sad.gif
20	6	happy	gif	subject06.happy.gif
21	8	rightlight	gif	subject08.rightlight.gif
22	13	sleepy	gif	subject13.sleepy.gif
23	3	leftlight	gif	subject03.leftlight.gif
24	5	sleepy	gif	subject05.sleepy.gif
25	10	centerlight	gif	subject10.centerlight.gif
26	8	normal	gif	subject08.normal.gif
27	1	gif	None	subject01.gif
28	7	leftlight	gif	subject07.leftlight.gif
29	11	happy	gif	subject11.happy.gif

```

In [ ]: x_train=df_train.loc[:,['category','filename']]
x_test=df_test.loc[:,['category','filename']]
y_train=df_train.loc[:,['subject']]
y_test=df_test.loc[:,['subject']]

```

```

In [ ]: y_train=y_train.to_numpy()
y_test=y_test.to_numpy()

```

```

In [ ]: y_train = y_train.tolist()
y_test = y_test.tolist()

```



```
In [ ]: detector = MTCNN()

def load_dataset1(dataset):
    faces = list()
    for filename in tqdm(dataset["filename"]):
        path = DIRECTORY_train + filename
        # get face
        face = extract_face_from_file(path)
        faces.append(face)
    return np.asarray(faces)
```

```
In [ ]: detector = MTCNN()

def load_dataset2(dataset):
    faces = list()
    for filename in tqdm(dataset["filename"]):
        path = DIRECTORY_test + filename
        # get face
        face = extract_face_from_file(path)
        faces.append(face)
    return np.asarray(faces)
```

```
In [ ]: x_test = load_dataset2(x_test)
x_train = load_dataset1(x_train)

print(x_test.shape)
print(x_train.shape)

100%|██████████| 30/30 [00:17<00:00, 1.67it/s]
100%|██████████| 135/135 [01:17<00:00, 1.74it/s]

(30, 160, 160)
(135, 160, 160)
```

3. Convolutional Neural Network Model

```
In [ ]: TRAINING_DATA_DIRECTORY = "data/train"
TESTING_DATA_DIRECTORY = "data/test"
NUM_CLASSES = 15
EPOCHS = 25
BATCH_SIZE = 20
NUMBER_OF_TRAINING_IMAGES = 135
NUMBER_OF_TESTING_IMAGES = 30
IMAGE_HEIGHT = 160
IMAGE_WIDTH = 160
```

```
In [ ]: import os

def save_keras_dataset(setname, dataset, labels, per_class):
    data = sorted(list(zip(labels, dataset)), key=lambda x: x[0])

    j = 0
    for label, gray_img in tqdm(data):
        j = (j% per_class) + 1

        directory = f"data/{setname}/class_{label}/"
        if not os.path.exists(directory):
            os.makedirs(directory)
        cv2.imwrite(f"{directory}class_{label}_{j}.png", gray_img)
```

```
In [ ]: import shutil
shutil.rmtree('data', ignore_errors=True)

# Save datasets
save_keras_dataset("test", x_test, y_test, 3)
save_keras_dataset("train", x_train, y_train, 8)

100%|██████████| 30/30 [00:00<00:00, 1084.57it/s]
100%|██████████| 135/135 [00:00<00:00, 1877.72it/s]
```

```
In [ ]: from tensorflow.keras.preprocessing.image import ImageDataGenerator
def data_generator():
    return ImageDataGenerator(
        rescale=1./255,

    )

def add_noise(img):
    """Add random noise to an image"""
    VARIABILITY = 35
    deviation = VARIABILITY*random.random()
    noise = np.random.normal(0, deviation, img.shape)
    img += noise
    np.clip(img, 0., 255.)
    return img
```

```
In [ ]: training_generator = data_generator().flow_from_directory(
    TRAINING_DATA_DIRECTORY,
    target_size=(IMAGE_WIDTH, IMAGE_HEIGHT),
    batch_size=BATCH_SIZE,
    class_mode='categorical',
    color_mode='grayscale'
)


testing_generator = data_generator().flow_from_directory(
    TESTING_DATA_DIRECTORY,
    target_size=(IMAGE_WIDTH, IMAGE_HEIGHT),
    class_mode='categorical',
    color_mode='grayscale'
)

validation_generator = data_generator().flow_from_directory(
    TESTING_DATA_DIRECTORY,
    target_size=(IMAGE_WIDTH, IMAGE_HEIGHT),
    class_mode='categorical',
    color_mode='grayscale',
    shuffle=False
)
```

Found 120 images belonging to 15 classes.
 Found 30 images belonging to 15 classes.
 Found 30 images belonging to 15 classes.

```
In [ ]: sample_images = testing_generator.next()[0]

f, xyarr = plt.subplots(3,3)
xyarr[0,0].imshow(sample_images[0])
xyarr[0,1].imshow(sample_images[1])
xyarr[0,2].imshow(sample_images[2])
xyarr[1,0].imshow(sample_images[3])
xyarr[1,1].imshow(sample_images[4])
xyarr[1,2].imshow(sample_images[5])
xyarr[2,0].imshow(sample_images[6])
xyarr[2,1].imshow(sample_images[7])
xyarr[2,2].imshow(sample_images[8])
plt.show()
```



```
In [ ]: import keras
class MCDropout(keras.layers.Dropout):
    def call(self, inputs):
        return super().call(inputs, training=True)
```

```
In [ ]: from tensorflow.keras import models
from tensorflow.keras.layers import Activation, ZeroPadding2D, MaxPooling2D, Conv2D, Flatten, Dense, Dropout
from tensorflow.keras import regularizers, constraints

model = models.Sequential()

model.add(Conv2D(32, kernel_size=(3, 3), activation='linear', input_shape=(IMAGE_WIDTH, IMAGE_HEIGHT, 1), padding='same'))
model.add(MaxPooling2D((2, 2)))

model.add(Conv2D(64, (3, 3), activation='relu', kernel_regularizer=regularizers.l2(l2=0.01)))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(128, (3, 3), activation='relu', kernel_regularizer=regularizers.l2(l2=0.01)))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())

model.add(Dense(512, activation='relu', kernel_initializer="glorot_uniform", kernel_regularizer=regularizers.l2(l2=0.01)))

model.add(MCDropout(rate=0.5))

model.add(Dense(NUM_CLASSES, activation='softmax', kernel_initializer="glorot_uniform"))
```

```
In [ ]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d_48 (Conv2D)	(None, 160, 160, 32)	320

max_pooling2d_24 (MaxPooling)	(None, 80, 80, 32)	0

conv2d_49 (Conv2D)	(None, 78, 78, 64)	18496

max_pooling2d_25 (MaxPooling)	(None, 39, 39, 64)	0

conv2d_50 (Conv2D)	(None, 37, 37, 128)	73856

max_pooling2d_26 (MaxPooling)	(None, 18, 18, 128)	0

flatten_8 (Flatten)	(None, 41472)	0

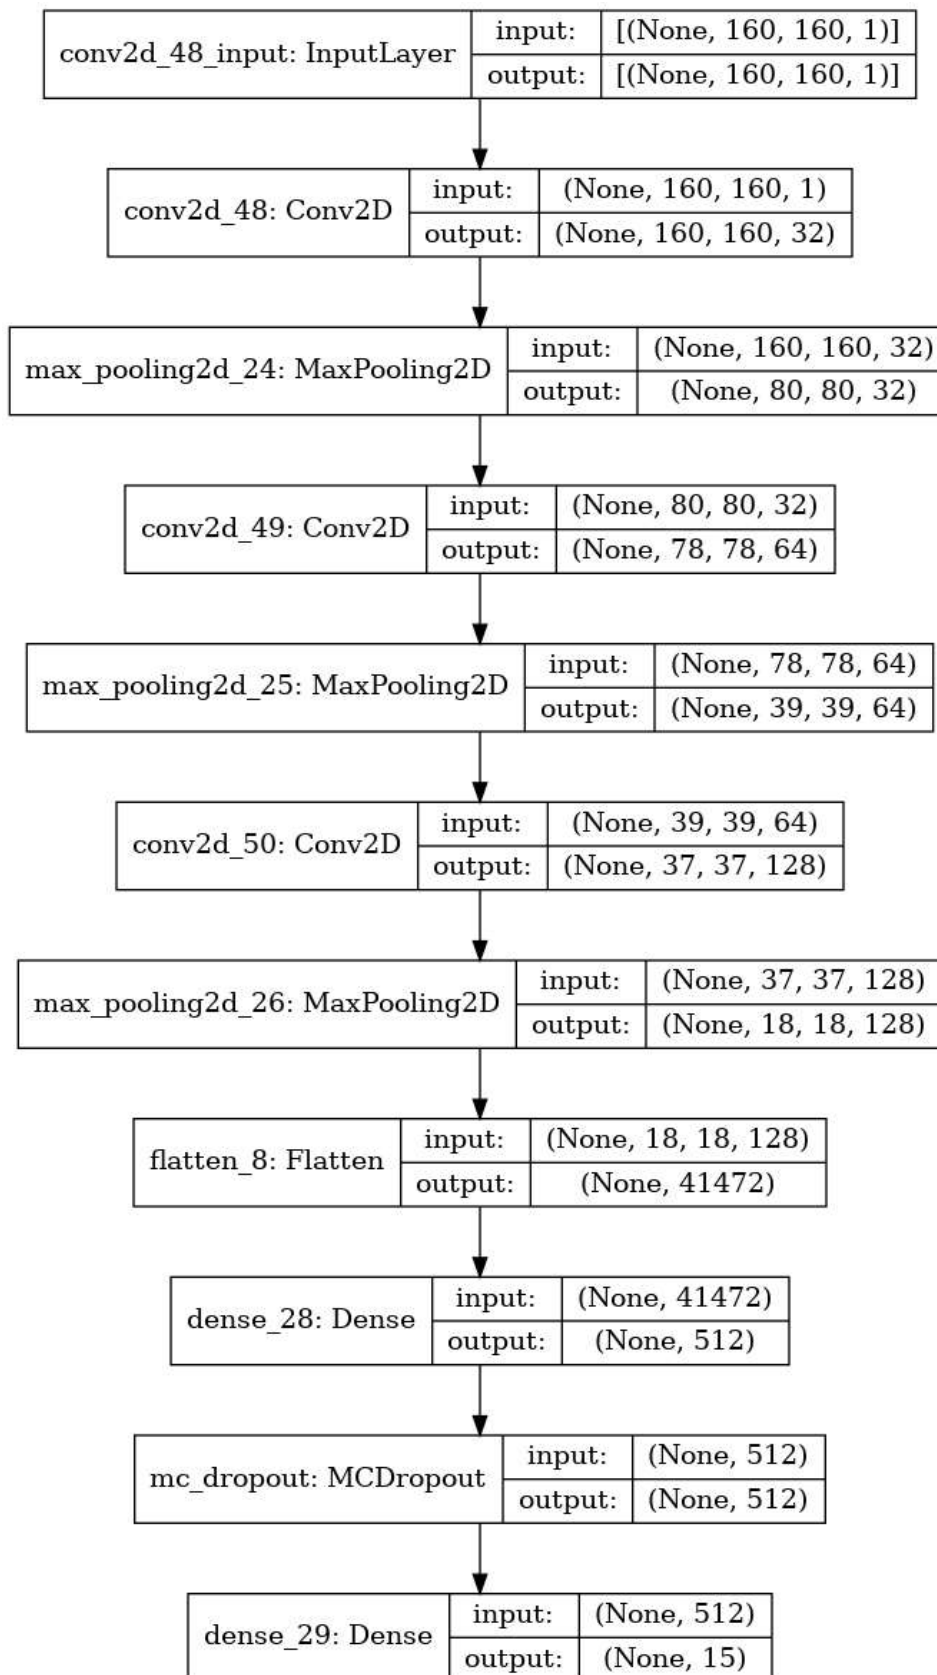
dense_28 (Dense)	(None, 512)	21234176

mc_dropout (MCDropout)	(None, 512)	0

dense_29 (Dense)	(None, 15)	7695
=====		
Total params: 21,334,543		
Trainable params: 21,334,543		
Non-trainable params: 0		

```
In [ ]: from keras.utils.vis_utils import plot_model
plot_model(model, show_shapes=True, show_layer_names=True)
```

Out[28]:



```
In [ ]: from tensorflow.keras import optimizers, losses
from tensorflow.keras.callbacks import EarlyStopping
early_stopping = EarlyStopping()

model.compile(
    loss=losses.CategoricalCrossentropy(from_logits=True),
    optimizer=optimizers.Adam(learning_rate=0.0003),
    metrics=["accuracy"]
)

history = model.fit(
    training_generator,
    steps_per_epoch=(NUMBER_OF_TRAINING_IMAGES//BATCH_SIZE ),
    epochs=EPOCHS,
    validation_data=testing_generator,
    shuffle=True,
    validation_steps=(NUMBER_OF_TESTING_IMAGES//BATCH_SIZE),
    #     callbacks=[early_stopping]
)
```


Epoch 1/25

/opt/conda/lib/python3.7/site-packages/keras/backend.py:4847: UserWarning: "`categorical_crossentropy` received `from_logits=True`, but the `output` argument was produced by a sigmoid or softmax activation and thus does not represent logits. Was this intended?"
 "`categorical_crossentropy` received `from_logits=True`, but "

6/6 [=====] - 4s 453ms/step - loss: 13.0665 - accuracy: 0.0583 - val_loss: 11.2362 - val_accuracy: 0.1000
Epoch 2/25
6/6 [=====] - 2s 402ms/step - loss: 10.2653 - accuracy: 0.1000 - val_loss: 8.9813 - val_accuracy: 0.1333
Epoch 3/25
6/6 [=====] - 2s 406ms/step - loss: 8.2131 - accuracy: 0.2250 - val_loss: 7.3162 - val_accuracy: 0.1667
Epoch 4/25
6/6 [=====] - 2s 415ms/step - loss: 6.6801 - accuracy: 0.2917 - val_loss: 6.0057 - val_accuracy: 0.3333
Epoch 5/25
6/6 [=====] - 3s 494ms/step - loss: 5.5678 - accuracy: 0.3250 - val_loss: 4.9027 - val_accuracy: 0.3000
Epoch 6/25
6/6 [=====] - 2s 405ms/step - loss: 4.6911 - accuracy: 0.3917 - val_loss: 4.3338 - val_accuracy: 0.4667
Epoch 7/25
6/6 [=====] - 2s 399ms/step - loss: 4.0437 - accuracy: 0.4917 - val_loss: 3.9233 - val_accuracy: 0.3667
Epoch 8/25
6/6 [=====] - 2s 408ms/step - loss: 3.4352 - accuracy: 0.6083 - val_loss: 3.6407 - val_accuracy: 0.4333
Epoch 9/25
6/6 [=====] - 2s 414ms/step - loss: 2.8870 - accuracy: 0.7583 - val_loss: 3.1343 - val_accuracy: 0.6000
Epoch 10/25
6/6 [=====] - 2s 402ms/step - loss: 2.6277 - accuracy: 0.7750 - val_loss: 2.8893 - val_accuracy: 0.6667
Epoch 11/25
6/6 [=====] - 2s 401ms/step - loss: 2.3975 - accuracy: 0.8333 - val_loss: 2.5442 - val_accuracy: 0.7667
Epoch 12/25
6/6 [=====] - 2s 412ms/step - loss: 2.1244 - accuracy: 0.9167 - val_loss: 2.5290 - val_accuracy: 0.8000
Epoch 13/25
6/6 [=====] - 2s 410ms/step - loss: 1.9782 - accuracy: 0.9250 - val_loss: 2.5517 - val_accuracy: 0.7000
Epoch 14/25
6/6 [=====] - 2s 402ms/step - loss: 1.8054 - accuracy: 0.9583 - val_loss: 2.5473 - val_accuracy: 0.7000
Epoch 15/25
6/6 [=====] - 2s 402ms/step - loss: 1.7525 - accuracy: 0.9750 - val_loss: 2.3258 - val_accuracy: 0.7333
Epoch 16/25
6/6 [=====] - 2s 399ms/step - loss: 1.6347 - accuracy: 0.9750 - val_loss: 1.9473 - val_accuracy: 0.8333
Epoch 17/25
6/6 [=====] - 2s 406ms/step - loss: 1.5042 - accuracy: 0.9917 - val_loss: 2.3997 - val_accuracy: 0.7333
Epoch 18/25
6/6 [=====] - 3s 485ms/step - loss: 1.4579 - accuracy: 0.9750 - val_loss: 1.8285 - val_accuracy: 0.8000
Epoch 19/25
6/6 [=====] - 2s 403ms/step - loss: 1.3716 - accuracy: 0.9833 - val_loss: 1.9576 - val_accuracy: 0.8333
Epoch 20/25
6/6 [=====] - 2s 402ms/step - loss: 1.3164 - accuracy: 0.9917 - val_loss: 1.7942 - val_accuracy: 0.8667
Epoch 21/25
6/6 [=====] - 2s 405ms/step - loss: 1.2431 - accuracy: 1.0000 - val_loss: 1.6066 - val_accuracy: 0.8333
Epoch 22/25
6/6 [=====] - 2s 405ms/step - loss: 1.1756 - accuracy: 1.0000 - val_loss: 1.8609 - val_accuracy: 0.7000
Epoch 23/25
6/6 [=====] - 2s 401ms/step - loss: 1.1020 - accuracy: 1.0000 - val_loss: 1.4665 - val_accuracy: 0.9000
Epoch 24/25
6/6 [=====] - 2s 412ms/step - loss: 1.0746 - accuracy: 0.9833 - val_loss: 1.5058 - val_accuracy: 0.8667
Epoch 25/25
6/6 [=====] - 2s 399ms/step - loss: 1.0788 - accuracy: 0.9750 - val_loss: 1.5114 - val_accuracy: 0.8667


```
In [ ]: plot_folder = "plot"
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0.1, 1])
plt.legend(loc='lower right')
```

Out[30]: <matplotlib.legend.Legend at 0x7f4e041a0790>



```
In [ ]: plot_folder = "plot"
plt.plot(history.history['loss'], label='loss')
plt.plot(history.history['val_loss'], label='val_loss')
plt.xlabel('Epoch')
plt.ylabel('Val Loss')
plt.legend(loc='lower right')
```

Out[31]: <matplotlib.legend.Legend at 0x7f4dff2e3c10>



```
In [ ]: from sklearn.metrics import classification_report, confusion_matrix, precision_recall_fscore_support
```

```
Y_pred = model.predict(validation_generator)
y_pred = np.argmax(Y_pred, axis=1)
print(classification_report(validation_generator.classes, y_pred))
print(validation_generator.classes)
print(y_pred)
print('Confusion Matrix')
print(confusion_matrix(validation_generator.classes, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	2
1	1.00	1.00	1.00	2
2	0.50	0.50	0.50	2
3	1.00	1.00	1.00	2
4	0.67	1.00	0.80	2
5	1.00	1.00	1.00	2
6	0.67	1.00	0.80	2
7	0.67	1.00	0.80	2
8	1.00	0.50	0.67	2
9	1.00	0.50	0.67	2
10	0.67	1.00	0.80	2
11	1.00	0.50	0.67	2
12	1.00	1.00	1.00	2
13	1.00	1.00	1.00	2
14	1.00	0.50	0.67	2
accuracy			0.83	30
macro avg	0.88	0.83	0.82	30
weighted avg	0.88	0.83	0.82	30

[0 0 1 1 2 2 3 3 4 4 5 5 6 6 7 7 8 8 9 9 10 10 11 11
12 12 13 13 14 14]

[0 0 1 1 2 10 3 3 4 4 5 5 6 6 7 7 2 8 9 4 10 10 6 11
12 12 13 13 14 7]

Confusion Matrix

```
[[2 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [0 2 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [0 0 1 0 0 0 0 0 0 0 1 0 0 0 0]
 [0 0 0 2 0 0 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 2 0 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 2 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 0 2 0 0 0 0 0 0 0 0]
 [0 0 1 0 0 0 0 0 1 0 0 0 0 0 0]
 [0 0 0 0 1 0 0 0 0 1 0 0 0 0 0]
 [0 0 0 0 0 0 0 0 0 0 2 0 0 0 0]
 [0 0 0 0 0 0 1 0 0 0 0 1 0 0 0]
 [0 0 0 0 0 0 0 0 0 0 0 0 2 0 0]
 [0 0 0 0 0 0 0 0 0 0 0 0 0 2 0]
 [0 0 0 0 0 0 0 0 1 0 0 0 0 0 1]]
```

```
In [ ]:
```